

Competitive Programming

Section 43 Editorial: String Algorithms I

Competitive Programming Course

2026-07-10

Table of contents

1	Editorial overview	3
1.1	Problem-specific proof map	3
1.2	Detailed solution guide	3
2	A. KMP Occurrences	5
2.1	Restatement	5
2.2	Observations and algorithm	5
2.3	Correctness proof	5
2.4	Complexity	5
2.5	Full C++ solution	5
2.6	Full Python solution	5
2.7	Tests that matter	6
3	B. All Borders	7
3.1	Restatement	7
3.2	Observations and algorithm	7
3.3	Correctness proof	7
3.4	Complexity	7
3.5	Full C++ solution	7
3.6	Full Python solution	7
3.7	Tests that matter	8
4	C. Minimal Period	9
4.1	Restatement	9
4.2	Observations and algorithm	9
4.3	Correctness proof	9
4.4	Complexity	9
4.5	Full C++ solution	9
4.6	Full Python solution	9
4.7	Tests that matter	10
5	What to carry forward	11
6	Testing discussion	12

1 Editorial overview

1.1 Problem-specific proof map

A. KMP Occurrences. The prefix function records the longest matched prefix after every text position; falling through failure links never discards a possible occurrence. Each character advances or falls only amortized constant times, so matching is $O(|\text{text}| + |\text{pattern}|)$.

B. All Borders. Repeatedly following `pi[length-1]` lists exactly the proper prefix lengths that are also suffixes, from longest to shortest; reverse the collected list for increasing output.

C. Minimal Period. If $p = n - \text{pi}[n-1]$ divides n , copies of the length- p prefix cover the string; otherwise no shorter border-compatible block can tile it, so the period is n .

1.2 Detailed solution guide

A. KMP Occurrences. Build the combined string `pattern + '#' + text`, where `#` is a separator that cannot appear in either input. The prefix function value `pi[i]` is the length of the longest prefix of the combined string ending at i . When `pi[i] == m`, an occurrence of the pattern has just ended. The end is at combined index i , the text starts at combined index $m + 1$, so the 1-based start position in the text is $i - 2*m + 1$. Overlapping matches are handled automatically because the prefix function falls back through borders instead of resetting to zero.

B. All Borders. A border of s is both a prefix and a suffix. The longest proper border has length `pi[n-1]`. If a string of length k is a border, then its own longest proper border has length `pi[k-1]`, and that is also a border of the original string. Repeating this prefix-link walk lists every border from longest to shortest; reversing gives increasing order. No border is skipped, because every shorter border of s must be a border of the longest border that contains it.

C. Minimal Period. If the longest border has length b , then the first candidate period is $p = n - b$: the suffix that repeats the prefix begins after p characters. This candidate is valid only when $n \% p == 0$; otherwise the repeated blocks do not tile the string and the answer is n . This is why strings such as `ababab` return 2, while `ababa` returns 5.

Common pitfalls. The separator must be absent from the alphabet, positions for A are 1-based, and B should not print the full length because a border is proper in this exercise.

The prefix function is the only real tool in this section. It is computed once in linear time and then interpreted in three ways.

- A builds `pattern + separator + text`; whenever `pi[i]` equals the pattern length, an occurrence ends at i .
- B follows prefix links from `pi[n-1]` to list every proper border of the full string.
- C uses the longest border to derive the candidate period $n - \text{pi}[n-1]$, then checks divisibility.

The important proof idea is that prefix links enumerate exactly the shorter borders, so no possible border or period is skipped.

2 A. KMP Occurrences

2.1 Restatement

Print the number and 1-based positions of all occurrences of pattern in text.

2.2 Observations and algorithm

Build `pattern + '#' + text` and its prefix function. If pattern length is `m`, `pi[i]==m` means a match ends at combined index `i`, whose 1-based text start is `i-2*m+1`. For `aba` in `abacaba`, this produces starts 1 and 5, including overlaps. Compute `pi` once, then scan its values.

2.3 Correctness proof

`pi[i]` is the longest pattern-prefix ending at `i`; consequently value `m` is exactly a complete occurrence. Failure links retain every shorter possible border, so no occurrence is skipped. QED.

2.4 Complexity

Time and memory are $O(|\text{pattern}| + |\text{text}|)$.

2.5 Full C++ solution

```
#include <bits/stdc++.h>
using namespace std; vector<int> pi(string s){vector<int> p(s.size());for(int i=1;i<(int)s.size();i++)p[i]=max(0,p[pi[i-1]]+s[i]==s[pi[i-1]]);return p;}
int main(){ios::sync_with_stdio(false);cin.tie(nullptr);string pat,text;if(!(cin>>pat>>text))return 0;vector<int> p=pi(pat+'#'+text);for(int i=0;i<p.size();i++)if(p[i]==pat.size())cout<<pat.size()<<" ";}
```

2.6 Full Python solution

```
import sys
def pi(s):
    p=[0]*len(s)
    for i in range(1,len(s)):
        j=p[i-1]
        while j and s[i]!=s[j]: j=p[j-1]
        p[i]=j+1
    return p
```

```

        if s[i]==s[j]: j+=1
        p[i]=j
    return p
def main():
    lines=sys.stdin.read().splitlines(); pat=lines[0]; text=lines[1]; s=pat+"#"+text; p=pi(s)
    for i,v in enumerate(p):
        if v==len(pat): ans.append(i-2*len(pat)+1)
    print(len(ans)); print(" ".join(map(str,ans)) if ans else "")
if __name__=="__main__": main()

```

2.7 Tests that matter

Test pattern longer than text, equal strings, no match, and overlapping **ana** in **anana**. The separator must be outside the input alphabet.

3 B. All Borders

3.1 Restatement

Print all proper border lengths of a string in increasing order.

3.2 Observations and algorithm

The longest proper border is $pi[n-1]$. Repeatedly set $k=pi[k-1]$ to follow the borders of that border; collect lengths and reverse for increasing output. For `ababa`, this visits 3 then 1.

3.3 Correctness proof

Every shorter border of the string is a border of its longest border, and every prefix link is a border by definition, so the walk lists exactly all borders. QED.

3.4 Complexity

Prefix construction and the link walk are $O(n)$ time and $O(n)$ memory.

3.5 Full C++ solution

```
#include <bits/stdc++.h>
using namespace std; int main(){ios::sync_with_stdio(false);cin.tie(nullptr);string s;if(!(C
```

3.6 Full Python solution

```
import sys
def pi(s):
    p=[0]*len(s)
    for i in range(1,len(s)):
        j=p[i-1]
        while j and s[i]!=s[j]: j=p[j-1]
        if s[i]==s[j]: j+=1
        p[i]=j
    return p
```

```
def main():
    s=sys.stdin.readline().strip(); p=pi(s); ans=[]; k=p[-1] if s else 0
    while k: ans.append(k); k=p[k-1]
    print(" ".join(map(str,ans[::-1])))
if __name__=="__main__": main()
```

3.7 Tests that matter

Test a singleton, `aaaa`, and a string without a border. Do not output `n`: this exercise asks only for proper borders.

4 C. Minimal Period

4.1 Restatement

Print the shortest period length if the string is repetitions of a block, otherwise n .

4.2 Observations and algorithm

Let $p = n - \text{pi}[n-1]$. It is a tiling period only if $n \% p == 0$; otherwise print n . For `abcabcabc`, p is 3; for `ababa`, candidate p is 2 but answer is 5.

4.3 Correctness proof

A border of length $n-p$ proves equality after shift p . Divisibility proves these shifts cover the whole string as equal blocks; without it, no shorter block tiles the string. QED.

4.4 Complexity

Time is $O(n)$ and memory is $O(n)$.

4.5 Full C++ solution

```
#include <bits/stdc++.h>
using namespace std; int main(){ios::sync_with_stdio(false);cin.tie(nullptr);string s;if(!(C
```

4.6 Full Python solution

```
import sys
def pi(s):
    p=[0]*len(s)
    for i in range(1,len(s)):
        j=p[i-1]
        while j and s[i]!=s[j]: j=p[j-1]
        if s[i]==s[j]: j+=1
        p[i]=j
    return p
```

```
def main():
    s=sys.stdin.readline().strip(); n=len(s); k=n-pi(s)[-1]
    print(k if n%k==0 else n)
if __name__=="__main__": main()
```

4.7 Tests that matter

Test one character, all one character, a prime-length nonperiodic string, and a near-repeat with a different final character.

5 What to carry forward

Advanced algorithms become manageable when you state the invariant first, then code only the operations needed to preserve it.

6 Testing discussion

For random strings of length at most 30, compare KMP positions with direct substring comparison, borders with a loop over lengths, and the period with a loop over every divisor of the length. Run:

```
CP_TARGET=solution python3 sections/43_string_algorithms_i/check.py
```